

- Logistic Regression as a Neural Network
- Python and Vectorisation
- Video: Vectorisation 3 min
- Video: More Vectorisation Examples 6 min
- Reading: Clarification of "it" 10 min
- Video: Vectorising Logistic Regression 7 min
- Video: Vectorising Logistic Regression: Gradients, Output 8 min
- Video: Broadcasting in Python 11 min
- Video: A note on performance vectors 6 min
- Video: Quick tour of Jupyter/Python Notebooks 3 min
- Video: Calculation of logistic regression cost function (optional) 7 min
- Practice Questions
- Quiz: Neural Network Basics 10 questions
- Programming Assignments
- Heroes of Deep Learning (optional)

Quiz - 30 min

Neural Network Basics

Submit your assignment

Due DATE Jan 16, 9:59 PM CST

Try again

Receive grade

To PASS 40% or higher

Grade 100%

View Feedback

We keep your highest score



✔ Congratulations! You passed!

To PASS 40% or higher

Keep Learning

GRADE 100%

Neural Network Basics

LEVEL: DISCUSSION GRADE
100%

1. What does a neuron compute? 1 / 1 point
- ☐ A neuron computes an activation function followed by a linear function ($z = Wx + b$)

☒ A neuron computes a linear function ($z = Wx + b$) followed by an activation function

☐ A neuron computes the mean of all features before applying the output to an activation function

☐ A neuron computes a function g that scales the input x linearly ($Wx + b$)
- ✔ Correct
- Correct, we generally say that the output of a neuron is $a = g(Wx + b)$ where g is the activation function (sigmoid, tanh, ReLU, ...).
2. Which of these is the "Logistic Loss"? 1 / 1 point
- ☐ $D^{(1)}(y^{(1)}, \hat{y}^{(1)}) = \max(0, y^{(1)} - \hat{y}^{(1)})$

☒ $D^{(1)}(y^{(1)}, \hat{y}^{(1)}) = -(y^{(1)} \log(\hat{y}^{(1)}) + (1 - y^{(1)}) \log(1 - \hat{y}^{(1)}))$

☐ $D^{(1)}(y^{(1)}, \hat{y}^{(1)}) = |y^{(1)} - \hat{y}^{(1)}|$

☐ $D^{(1)}(y^{(1)}, \hat{y}^{(1)}) = |y^{(1)} - \hat{y}^{(1)}|^2$
- ✔ Correct
- Correct, this is the logistic loss you've been in lecture!
3. Suppose `img` is a (32,32,3) array, representing a 32x32 image with 3 color channels: red, green and blue. How do you reshape this into a column vector? 1 / 1 point
- ☐ `x = img.reshape(32*32,3)`

☐ `x = img.reshape(3,32*32)`

☐ `x = img.reshape(3,32*32*3)`

☒ `x = img.reshape(32*32*3,1)`
- ✔ Correct
4. Consider the two following random arrays "a" and "b": 1 / 1 point
- ```
1 a = np.random.randn(2, 3) # a.shape = (2, 3)
2 b = np.random.randn(2, 3) # b.shape = (2, 3)
3 c = a + b
```
- What will be the shape of "c"?
- ☐ The computation cannot happen because the sizes don't match. It's going to be "Error"

☐ `c.shape = (3, 2)`

☐ `c.shape = (2, 1)`

☒ `c.shape = (2, 3)`
- ✔ Correct
- Yes! This is broadcasting. `b` (column vector) is copied 3 times so that it can be summed to each column of `a`.
5. Consider the two following random arrays "a" and "b": 1 / 1 point
- ```
1 a = np.random.randn(4, 3) # a.shape = (4, 3)
2 b = np.random.randn(2, 4) # b.shape = (2, 4)
3 c = a*b
```
- What will be the shape of "c"?
- ☐ `c.shape = (4, 3)`

☐ `c.shape = (3, 3)`

☐ `c.shape = (4,2)`

☒ The computation cannot happen because the sizes don't match. It's going to be "Error"
- ✔ Correct
- Indeed! In numpy the "*" operator indicates element-wise multiplication. It is different from "np.dot()". If you would try `c = np.dot(a,b)` you would get `c.shape = (4, 2)`.
6. Suppose you have n_x input features per example. Recall that $X = [x^{(1)} x^{(2)} \dots x^{(m)}]$. What is the dimension of X ? 1 / 1 point
- ☐ $(m, 1)$

☒ (n_x, m)

☐ $(1, m)$

☐ (m, n_x)
- ✔ Correct
7. Recall that "np.dot(a,b)" performs a matrix multiplication on `a` and `b`, whereas "a*b" performs an element-wise multiplication. 1 / 1 point
- Consider the two following random arrays "a" and "b":
- ```
1 a = np.random.randn(12288, 128) # a.shape = (12288, 128)
2 b = np.random.randn(128, 45) # b.shape = (128, 45)
3 c = np.dot(a,b)
```
- What is the shape of `c`?
- ☐ `c.shape = (12288, 150)`

☐ The computation cannot happen because the sizes don't match. It's going to be "Error"

☒ `c.shape = (12288, 45)`

☐ `c.shape = (150, 150)`
- ✔ Correct
- Correct, remember that a `np.dot(a, b)` has shape (number of rows of `a`, number of columns of `b`). The sizes match because: "number of columns of `a` = 128 = number of rows of `b`".
8. Consider the following code snippet: 1 / 1 point
- ```
1 # a.shape = (3,4)
2 # b.shape = (4,3)
3
4 for j in range(3):
5     for i in range(4):
6         c[i][j] = a[i][j] + b[j]
```
- How do you vectorise this?
- ☐ `c = a.T + b.T`

☐ `c = b + b`

☐ `c = a.T + b`

☒ `c = a + b.T`
- ✔ Correct
9. Consider the following code: 1 / 1 point
- ```
1 # a = np.random.randn(3, 3)
2 # b = np.random.randn(3, 3)
3 c = a*b
```
- What will be `c` if you're not sure, feel free to run this in python to find out.
- ☒ This will involve broadcasting, so `b` is copied three times to become (3,3), and `*` is an element-wise product so `c.shape` will be (3,3)

☐ This will involve broadcasting, so `b` is copied three times to become (3, 3), and `*` invokes a matrix multiplication operation of two 3x3 matrices so `c.shape` will be (3, 3)

☐ This will multiply a 3x3 matrix `a` with a 3x1 vector, thus resulting in a 3x1 vector. That is, `c.shape = (3, 1)`.

☐ It will lead to an error since you cannot use "\*" to operate on these two matrices. You need to instead use `np.dot(a,b)`
- ✔ Correct
10. Consider the following computation graph. 1 / 1 point
- a

b

c

U = a<sup>T</sup> b

V = a<sup>T</sup> c

W = b<sup>T</sup> c

J = U + V + W
- What is the output `J`?
- ☐ `J = (c - 1)T b + a`

☒ `J = (a - 1)T b + c`

☐ `J = aT b + bT c + aT c`

☐ `J = bT - 1)T b + a`
- ✔ Correct
- Yes! `J = U + V + W = aT b + aT c + bT c = (a + b)T c = (a + b)T (a - 1)T b + c`.